

# ISLR Lab 3.6

Neev Shaw

September 12, 2025

```
[1]: library(MASS)  
library(ISLR)
```

```
[2]: fix(Boston)  
print(names(Boston))
```

```
[1] "crim"      "zn"        "indus"     "chas"      "nox"       "rm"        "age"  
[8] "dis"       "rad"       "tax"       "ptratio"   "black"     "lstat"     "medv"
```

```
[3]: lm.fit=lm(medv~lstat)
```

```
Error in eval(predvars, data, env): object 'medv' not found  
Traceback:
```

```
1. eval(mf, parent.frame())  
2. eval(mf, parent.frame())  
3. stats::model.frame(formula = medv ~ lstat, drop.unused.levels = TRUE)  
4. model.frame.default(formula = medv ~ lstat, drop.unused.levels = TRUE)  
5. eval(predvars, data, env)  
6. eval(predvars, data, env)  
7. .handleSimpleError(function (cnd)  
{  
  .  
  watcher$capture_plot_and_output()  
  cnd <- sanitize_call(cnd)  
  watcher$push(cnd)  
  switch(on_error, continue = invokeRestart("eval_continue"),  
        stop = invokeRestart("eval_stop"), error = NULL)  
}, "object 'medv' not found", base::quote(eval(predvars, data,  
env)))
```

```
[4]: lm.fit=lm(medv~lstat, data=Boston)  
attach(Boston)  
lm.fit=lm(medv~lstat)
```

```
[5]: lm.fit
```

```
Call:
lm(formula = medv ~ lstat)
```

```
Coefficients:
(Intercept)      lstat
            34.55     -0.95
```

[6]: `summary(lm.fit)`

```
Call:
lm(formula = medv ~ lstat)
```

```
Residuals:
    Min      1Q  Median      3Q      Max 
-15.168  -3.990  -1.318   2.034  24.500
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 34.55384    0.56263   61.41 <2e-16 ***
lstat       -0.95005    0.03873  -24.53 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 6.216 on 504 degrees of freedom
Multiple R-squared:  0.5441,    Adjusted R-squared:  0.5432 
F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16
```

[7]: `print(names(lm.fit))`

```
[1] "coefficients"   "residuals"        "effects"        "rank"      
[5] "fitted.values"  "assign"          "qr"             "df.residual"  
[9] "xlevels"         "call"           "terms"          "model"
```

[8]: `print(coef(lm.fit))`

```
(Intercept)      lstat
34.5538409  -0.9500494
```

[9]: `print(confint(lm.fit))`

```
              2.5 %      97.5 %    
(Intercept) 33.448457 35.6592247
lstat       -1.026148 -0.8739505
```


```
[10]: print(predict(lm.fit, data.frame(lstat=c(5,10 ,15)), interval ="confidence"))
```

	fit	lwr	upr
1	29.80359	29.00741	30.59978
2	25.05335	24.47413	25.63256
3	20.30310	19.73159	20.87461


```
[11]: print(predict(lm.fit, data.frame(lstat=c(5,10 ,15)), interval ="prediction"))
```

	fit	lwr	upr
1	29.80359	17.565675	42.04151
2	25.05335	12.827626	37.27907
3	20.30310	8.077742	32.52846


```
[12]: plot(lstat, medv)
abline(lm.fit)
```




```
[13]: plot(lstat, medv)
abline(lm.fit, lwd=3)
```




```
[14]: plot(lstat, medv)
abline(lm.fit, lwd=3, col="red")
```




```
[15]: plot(lstat ,medv ,col="red")
```




```
[16]: plot(lstat ,medv ,pch =20)
```




```
[17]: plot(lstat ,medv ,pch ="+")
```




```
[18]: plot(1:20,1:20,pch =1:20)
```




```
[19]: par(mfrow=c(2,2))
plot(lm.fit)
```




```
[20]: plot(predict(lm.fit), residuals(lm.fit))
```



```
[21]: plot(predict(lm.fit), rstudent(lm.fit))
```



```
[22]: plot(hatvalues(lm.fit))
```



```
[23]: which.max(hatvalues(lm.fit))
```

375: 375

```
[24]: lm.fit=lm(medv~lstat+age ,data=Boston )
summary(lm.fit)
```

Call:  
lm(formula = medv ~ lstat + age, data = Boston)

Residuals:  
Min 1Q Median 3Q Max  
-15.981 -3.978 -1.283 1.968 23.158

```

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 33.22276   0.73085 45.458 < 2e-16 ***
lstat       -1.03207   0.04819 -21.416 < 2e-16 ***
age         0.03454   0.01223   2.826 0.00491 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 6.173 on 503 degrees of freedom  
 Multiple R-squared: 0.5513, Adjusted R-squared: 0.5495  
 F-statistic: 309 on 2 and 503 DF, p-value: < 2.2e-16

```
[25]: lm.fit=lm(medv~., data=Boston )
summary(lm.fit)
```

```

Call:
lm(formula = medv ~ ., data = Boston)

Residuals:
```

Min	1Q	Median	3Q	Max
-15.595	-2.730	-0.518	1.777	26.199

Coefficients:

```

            Estimate Std. Error t value Pr(>|t|)
(Intercept) 3.646e+01 5.103e+00 7.144 3.28e-12 ***
crim        -1.080e-01 3.286e-02 -3.287 0.001087 **
zn          4.642e-02 1.373e-02  3.382 0.000778 ***
indus       2.056e-02 6.150e-02  0.334 0.738288
chas        2.687e+00 8.616e-01  3.118 0.001925 **
nox        -1.777e+01 3.820e+00 -4.651 4.25e-06 ***
rm          3.810e+00 4.179e-01  9.116 < 2e-16 ***
age        6.922e-04 1.321e-02  0.052 0.958229
dis        -1.476e+00 1.995e-01 -7.398 6.01e-13 ***
rad         3.060e-01 6.635e-02  4.613 5.07e-06 ***
tax        -1.233e-02 3.760e-03 -3.280 0.001112 **
ptratio    -9.527e-01 1.308e-01 -7.283 1.31e-12 ***
black      9.312e-03 2.686e-03  3.467 0.000573 ***
lstat      -5.248e-01 5.072e-02 -10.347 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 4.745 on 492 degrees of freedom  
 Multiple R-squared: 0.7406, Adjusted R-squared: 0.7338  
 F-statistic: 108.1 on 13 and 492 DF, p-value: < 2.2e-16

```
[26]: library(car)
print(vif(lm.fit))
```

Loading required package: carData

	crim	zn	indus	chas	nox	rm	age	dis
1.792192	2.298758	3.991596	1.073995	4.393720	1.933744	3.100826	3.955945	
rad	tax	ptratio	black	lstat				
7.484496	9.008554	1.799084	1.348521	2.941491				

```
[27]: lm.fit1=lm(medv~.-age ,data=Boston )
summary (lm.fit1)
```

Call:  
`lm(formula = medv ~ . - age, data = Boston)`

Residuals:

Min	1Q	Median	3Q	Max
-15.6054	-2.7313	-0.5188	1.7601	26.2243

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	36.436927	5.080119	7.172	2.72e-12 ***
crim	-0.108006	0.032832	-3.290	0.001075 **
zn	0.046334	0.013613	3.404	0.000719 ***
indus	0.020562	0.061433	0.335	0.737989
chas	2.689026	0.859598	3.128	0.001863 **
nox	-17.713540	3.679308	-4.814	1.97e-06 ***
rm	3.814394	0.408480	9.338	< 2e-16 ***
dis	-1.478612	0.190611	-7.757	5.03e-14 ***
rad	0.305786	0.066089	4.627	4.75e-06 ***
tax	-0.012329	0.003755	-3.283	0.001099 **
ptratio	-0.952211	0.130294	-7.308	1.10e-12 ***
black	0.009321	0.002678	3.481	0.000544 ***
lstat	-0.523852	0.047625	-10.999	< 2e-16 ***
---				
Signif. codes:	0 ***	0.001 **	0.01 *	0.05 .
	''	'	'	'
	'	'	'	'

Residual standard error: 4.74 on 493 degrees of freedom

Multiple R-squared: 0.7406, Adjusted R-squared: 0.7343

F-statistic: 117.3 on 12 and 493 DF, p-value: < 2.2e-16

```
[28]: lm.fit1=update(lm.fit , ~.-age)
```

```
[29]: summary (lm(medv~lstat*age ,data=Boston))
```

```

Call:
lm(formula = medv ~ lstat * age, data = Boston)

Residuals:
    Min      1Q  Median      3Q     Max 
-15.806 -4.045 -1.333  2.085 27.552 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 36.0885359  1.4698355  24.553 < 2e-16 ***
lstat        -1.3921168  0.1674555  -8.313 8.78e-16 ***
age          -0.0007209  0.0198792  -0.036  0.9711    
lstat:age     0.0041560  0.0018518   2.244  0.0252 *  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.149 on 502 degrees of freedom
Multiple R-squared:  0.5557,    Adjusted R-squared:  0.5531 
F-statistic: 209.3 on 3 and 502 DF,  p-value: < 2.2e-16

```

[30]: `lm.fit2=lm(medv~lstat+I(lstat^2))  
summary (lm.fit2)`

```

Call:
lm(formula = medv ~ lstat + I(lstat^2))

Residuals:
    Min      1Q  Median      3Q     Max 
-15.2834 -3.8313 -0.5295  2.3095 25.4148 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 42.862007  0.872084  49.15 <2e-16 ***
lstat        -2.332821  0.123803 -18.84 <2e-16 ***
I(lstat^2)   0.043547  0.003745  11.63 <2e-16 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.524 on 503 degrees of freedom
Multiple R-squared:  0.6407,    Adjusted R-squared:  0.6393 
F-statistic: 448.5 on 2 and 503 DF,  p-value: < 2.2e-16

```

[31]: `lm.fit=lm(medv~lstat)  
print(anova(lm.fit ,lm.fit2))`


## Analysis of Variance Table

```

Model 1: medv ~ lstat
Model 2: medv ~ lstat + I(lstat^2)
  Res.Df   RSS Df Sum of Sq    F    Pr(>F)
1     504 19472
2     503 15347  1     4125.1 135.2 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
[32]: par(mfrow=c(2,2))
plot(lm.fit2)
```



[33]: `lm.fit5=lm(medv~poly(lstat ,5))  
summary(lm.fit5)`

```
Call:  
lm(formula = medv ~ poly(lstat, 5))  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-13.5433 -3.1039 -0.7052  2.0844 27.1153  
  
Coefficients:  
            Estimate Std. Error t value Pr(>|t|)  
(Intercept) 22.5328    0.2318  97.197 < 2e-16 ***  
poly(lstat, 5)1 -152.4595   5.2148 -29.236 < 2e-16 ***  
poly(lstat, 5)2   64.2272   5.2148 12.316 < 2e-16 ***  
poly(lstat, 5)3  -27.0511   5.2148 -5.187 3.10e-07 ***  
poly(lstat, 5)4   25.4517   5.2148  4.881 1.42e-06 ***  
poly(lstat, 5)5  -19.2524   5.2148 -3.692 0.000247 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 5.215 on 500 degrees of freedom  
Multiple R-squared:  0.6817,    Adjusted R-squared:  0.6785  
F-statistic: 214.2 on 5 and 500 DF,  p-value: < 2.2e-16
```

[34]: `summary(lm(medv~log(rm) ,data=Boston))`

```
Call:  
lm(formula = medv ~ log(rm), data = Boston)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-19.487 -2.875 -0.104   2.837 39.816  
  
Coefficients:  
            Estimate Std. Error t value Pr(>|t|)  
(Intercept) -76.488      5.028  -15.21  <2e-16 ***  
log(rm)       54.055     2.739   19.73  <2e-16 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 6.915 on 504 degrees of freedom  
Multiple R-squared:  0.4358,    Adjusted R-squared:  0.4347  
F-statistic: 389.3 on 1 and 504 DF,  p-value: < 2.2e-16
```

```
[35]: fix(Carseats )
print(names(Carseats ))
```

```
[1] "Sales"          "CompPrice"      "Income"        "Advertising"    "Population"
[6] "Price"          "ShelveLoc"      "Age"           "Education"     "Urban"
[11] "US"
```

```
[36]: lm.fit=lm(Sales~. + Income : Advertising + Price:Age ,data=Carseats )
summary(lm.fit)
```

```
Call:
lm(formula = Sales ~ . + Income:Advertising + Price:Age, data = Carseats)

Residuals:
    Min      1Q  Median      3Q      Max 
-2.9208 -0.7503  0.0177  0.6754  3.3413 

Coefficients:
              Estimate Std. Error t value Pr(>|t|)    
(Intercept) 6.5755654  1.0087470  6.519 2.22e-10 ***
CompPrice    0.0929371  0.0041183 22.567 < 2e-16 ***
Income       0.0108940  0.0026044  4.183 3.57e-05 ***
Advertising  0.0702462  0.0226091  3.107 0.002030 ** 
Population   0.0001592  0.0003679  0.433 0.665330  
Price        -0.1008064  0.0074399 -13.549 < 2e-16 ***
ShelveLocGood 4.8486762  0.1528378 31.724 < 2e-16 ***
ShelveLocMedium 1.9532620  0.1257682 15.531 < 2e-16 ***
Age          -0.0579466  0.0159506 -3.633 0.000318 *** 
Education    -0.0208525  0.0196131 -1.063 0.288361  
UrbanYes     0.1401597  0.1124019  1.247 0.213171  
USYes        -0.1575571  0.1489234 -1.058 0.290729  
Income:Advertising 0.0007510  0.0002784  2.698 0.007290 ** 
Price:Age     0.0001068  0.0001333  0.801 0.423812  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.011 on 386 degrees of freedom
Multiple R-squared:  0.8761,    Adjusted R-squared:  0.8719 
F-statistic:  210 on 13 and 386 DF,  p-value: < 2.2e-16
```

```
[37]: attach(Carseats )
print(contrasts (ShelveLoc ))
```

	Good	Medium
Bad	0	0
Good	1	0

Medium 0 1

```
[2]: LoadLibraries = function() {  
    library(ISLR)  
    library(MASS)  
    print("The libraries have been loaded.")  
}
```

```
[3]: print(LoadLibraries)
```

```
function ()  
{  
    library(ISLR)  
    library(MASS)  
    print("The libraries have been loaded.")  
}
```

```
[40]: LoadLibraries()
```

```
[1] "The libraries have been loaded."
```