

Assignment 5

Neev Shaw

October 3, 2025

Exercise 6

Part (a)

From the problem we know that

$$p(X) = \frac{e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2}}{1 + e^{\beta_0 + \beta_1 X_1 + \beta_2 X_2}}.$$

So we can plug in our approximations for each coefficient and find the probability for $X_1 = 40$ and $X_2 = 3.5$:

$$p(X) \approx \frac{e^{-6+0.05X_1+X_2}}{1 + e^{-6+0.05X_1+X_2}} = \boxed{0.378}$$

Part (b)

For $p(X) = 1/2$, we can see that the e^- term must equal 1, so the exponent $\beta_0 + \beta_1 X_1 + \beta_2 X_2$ must equal 0. Plugging in the value for each of the coefficients and X_2 then solving for X_1 gives:

$$-6 + 0.05X_1 + 3.5 = 0$$

$$X_1 = \boxed{50 \text{ hrs}}$$

Exercise 8

With $K = 1$, the training error rate is 0% because the nearest neighbor is simply the point itself, so each training data point is assigned the label it actually is. Since the average error rate is 18% this means the test error rate is $\frac{0\% + \text{Test}}{2} = 18\% \implies \text{Test} = 2 \cdot 18\% = 36\%$. 36% is higher than the logistic regression test error rate of 30% so we should prefer logistic regression.

Exercise 9

Part (a)

This means that $\frac{p(X)}{1 - p(X)} = 0.37$ so we solve for $p(X)$ to get that 37/137 of people with an odds of 0.37 will actually default.

Part (b)

We simply just plug 0.16 into the above formula to get that the odds are

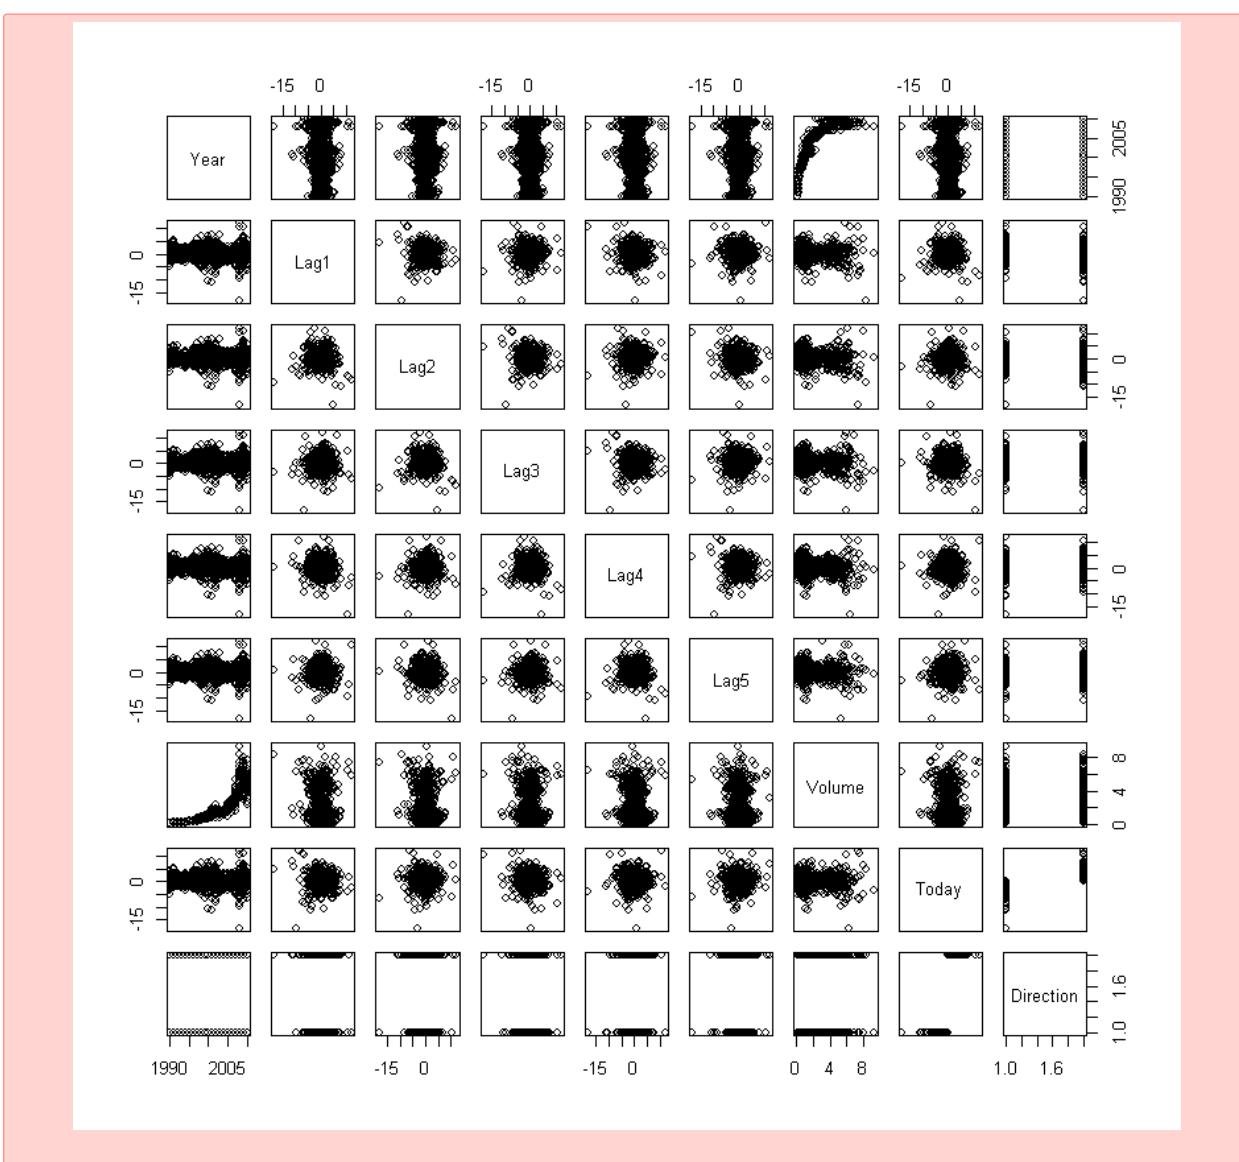
$$\frac{0.16}{1 - 0.16} \approx 0.19.$$

Exercise 10**Part (a)**

```
[1]: library(ISLR)
```

```
[2]: attach(Weekly)
fix(Weekly)
```

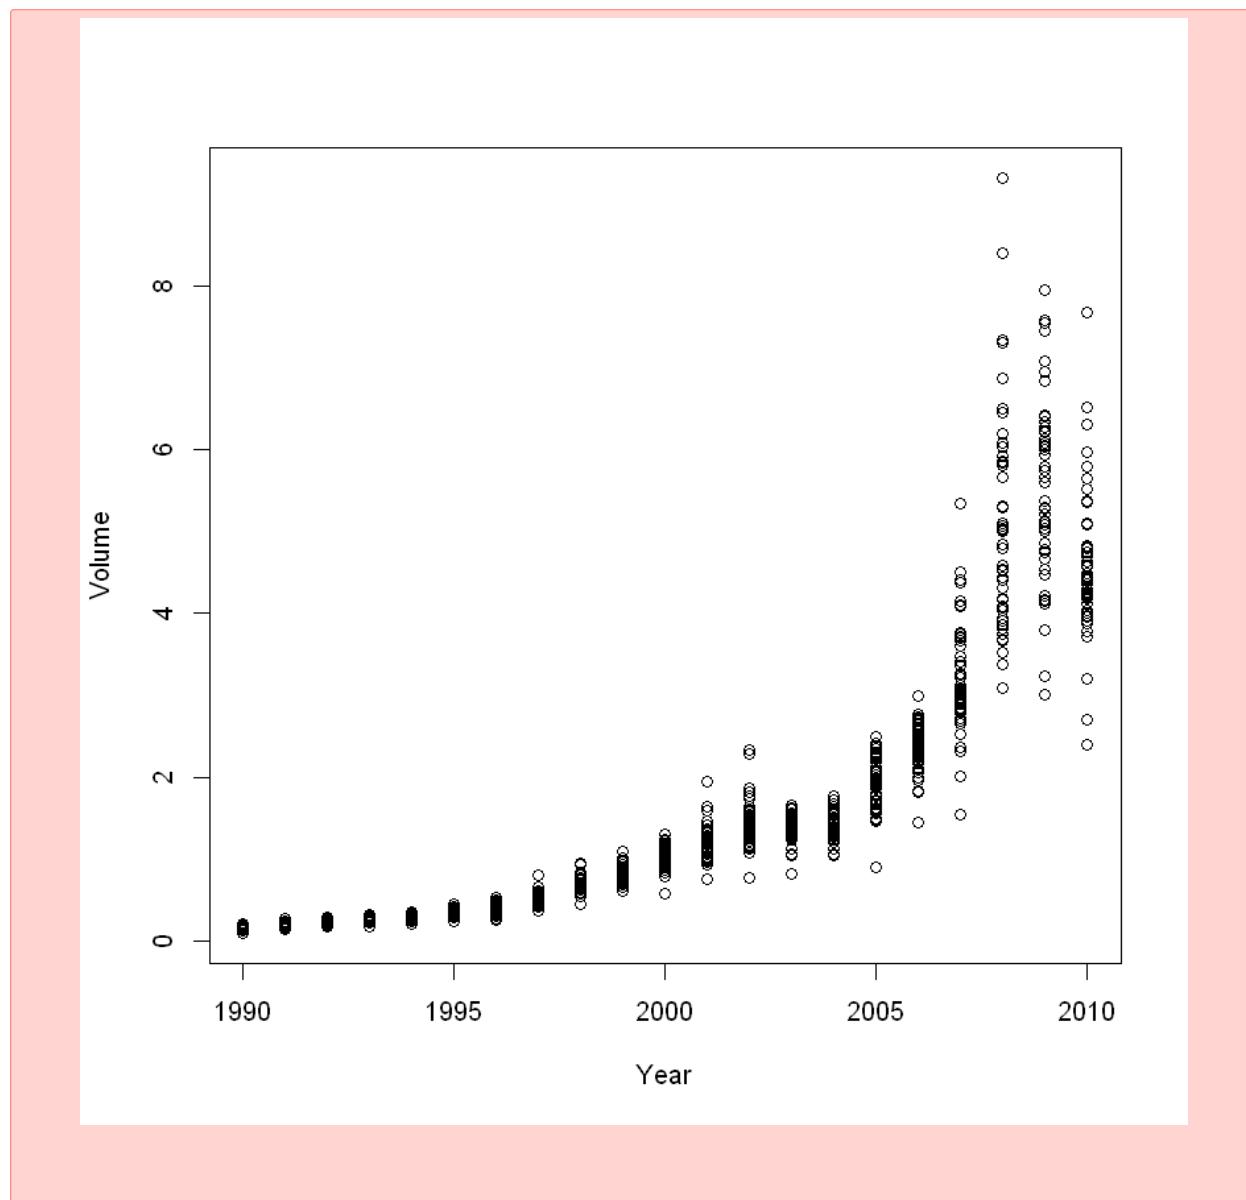
```
[3]: pairs(Weekly)
```



```
[4]: summary(Weekly)
```

Year	Lag1	Lag2	Lag3
Min. :1990	Min. :-18.1950	Min. :-18.1950	Min. :-18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260
Lag4	Lag5	Volume	Today
Min. :-18.1950	Min. :-18.1950	Min. :0.08747	Min. :-18.1950
1st Qu.: -1.1580	1st Qu.: -1.1660	1st Qu.:0.33202	1st Qu.: -1.1540
Median : 0.2380	Median : 0.2340	Median :1.00268	Median : 0.2410
Mean : 0.1458	Mean : 0.1399	Mean :1.57462	Mean : 0.1499
3rd Qu.: 1.4090	3rd Qu.: 1.4050	3rd Qu.:2.05373	3rd Qu.: 1.4050
Max. : 12.0260	Max. : 12.0260	Max. :9.32821	Max. : 12.0260
Direction			
Down:484			
Up :605			

```
[7]: plot(Year, Volume)
```



The only really noticeable trend is that volume seems to increase (almost exponentially?) with year.

Part (b)

```
[11]: glm.fits =glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume, data=Weekly  
,family=binomial )  
summary(glm.fits)
```

```
Call:  
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +  
Volume, family = binomial, data = Weekly)
```

```
Coefficients:
```

```

Estimate Std. Error z value Pr(>|z|)
(Intercept) 0.26686 0.08593 3.106 0.0019 **
Lag1        -0.04127 0.02641 -1.563 0.1181
Lag2         0.05844 0.02686 2.175 0.0296 *
Lag3        -0.01606 0.02666 -0.602 0.5469
Lag4        -0.02779 0.02646 -1.050 0.2937
Lag5        -0.01447 0.02638 -0.549 0.5833
Volume      -0.02274 0.03690 -0.616 0.5377
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom
Residual deviance: 1486.4 on 1082 degrees of freedom
AIC: 1500.4

Number of Fisher Scoring iterations: 4

```

The only predictor that is statistically significant is Lag2, but even that is only significant to the 0.05 level.

Part (c)

```
[14]: glm.probs=predict(glm.fits,type="response")
glm.pred=rep("Down" ,1089)
glm.pred[glm.probs >.5]="Up"
print(table(glm.pred ,Direction ))
```

Direction	
glm.pred	Down Up
Down	54 48
Up	430 557

```
[15]: (557+54)/1089
```

0.561065197428834

The model has an accuracy of 56.1%. The confusion matrix tells us what the model predicted compared to the actual answer. By looking at the off-diagonal, we can see that the model tends to predict up when the direction is actually down, more often than it predicts down when the direction is up.

Part (d)

```
[21]: train=(Year<2009)
Weekly.2009= Weekly[!train ,]
Direction.2009= Direction[!train]

[22]: glm.fits =glm(Direction~Lag2, data=Weekly ,family=binomial, subset=train)

[23]: glm.probs=predict (glm.fits,Weekly.2009, type="response")
glm.pred=rep("Down",length(glm.probs))
glm.pred[glm.probs >.5]="Up"
table(glm.pred ,Direction.2009)
```

Direction.2009		
glm.pred	Down	Up
Down	9	5
Up	34	56

```
[30]: mean(glm.pred==Direction.2009)
```

0.625

The confusion matrix is given above and the fraction of correct predictions is 5/8.