

Project 5

Neev Shaw

November 9, 2025

Part 1: Build an initial model

```
[1]: train = read.csv("TitanicTrain.csv")
train <- na.omit(train)
fix(train)
```

```
[2]: summary(train)
```

PassengerId	Survived	Pclass	Name
Min. : 1.0	Min. :0.0000	Min. :1.000	Length:714
1st Qu.:222.2	1st Qu.:0.0000	1st Qu.:1.000	Class :character
Median :445.0	Median :0.0000	Median :2.000	Mode :character
Mean :448.6	Mean :0.4062	Mean :2.237	
3rd Qu.:677.8	3rd Qu.:1.0000	3rd Qu.:3.000	
Max. :891.0	Max. :1.0000	Max. :3.000	
Sex	Age	SibSp	Parch
Length:714	Min. : 0.42	Min. :0.0000	Min. :0.0000
Class :character	1st Qu.:20.12	1st Qu.:0.0000	1st Qu.:0.0000
Mode :character	Median :28.00	Median :0.0000	Median :0.0000
	Mean :29.70	Mean :0.5126	Mean :0.4314
	3rd Qu.:38.00	3rd Qu.:1.0000	3rd Qu.:1.0000
	Max. :80.00	Max. :5.0000	Max. :6.0000
Ticket	Fare	Cabin	Embarked
Length:714	Min. : 0.00	Length:714	Length:714
Class :character	1st Qu.: 8.05	Class :character	Class :character
Mode :character	Median : 15.74	Mode :character	Mode :character
	Mean : 34.69		
	3rd Qu.: 33.38		
	Max. :512.33		

```
[3]: train$Sex = as.factor(train$Sex)
train$Pclass = as.factor(train$Pclass)
```

```
[4]: model = glm(Survived~Sex+Age+Pclass, family="binomial", data=train)
summary(model)
```

Call:

```

glm(formula = Survived ~ Sex + Age + Pclass, family = "binomial",
    data = train)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 3.777013  0.401123  9.416 < 2e-16 ***
Sexmale     -2.522781  0.207391 -12.164 < 2e-16 ***
Age        -0.036985  0.007656 -4.831 1.36e-06 ***
Pclass2     -1.309799  0.278066 -4.710 2.47e-06 ***
Pclass3     -2.580625  0.281442 -9.169 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 964.52 on 713 degrees of freedom
Residual deviance: 647.28 on 709 degrees of freedom
AIC: 657.28

Number of Fisher Scoring iterations: 5

```

```
[5]: glm.probs=predict(model,type="response")
glm.pred=rep(0,nrow(train))
glm.pred[glm.probs >0.5]=1
table(glm.pred ,train$Survived)
mean(glm.pred==train$Survived)
```

```

glm.pred 0 1
 0 356 83
 1 68 207
0.788515406162465

```

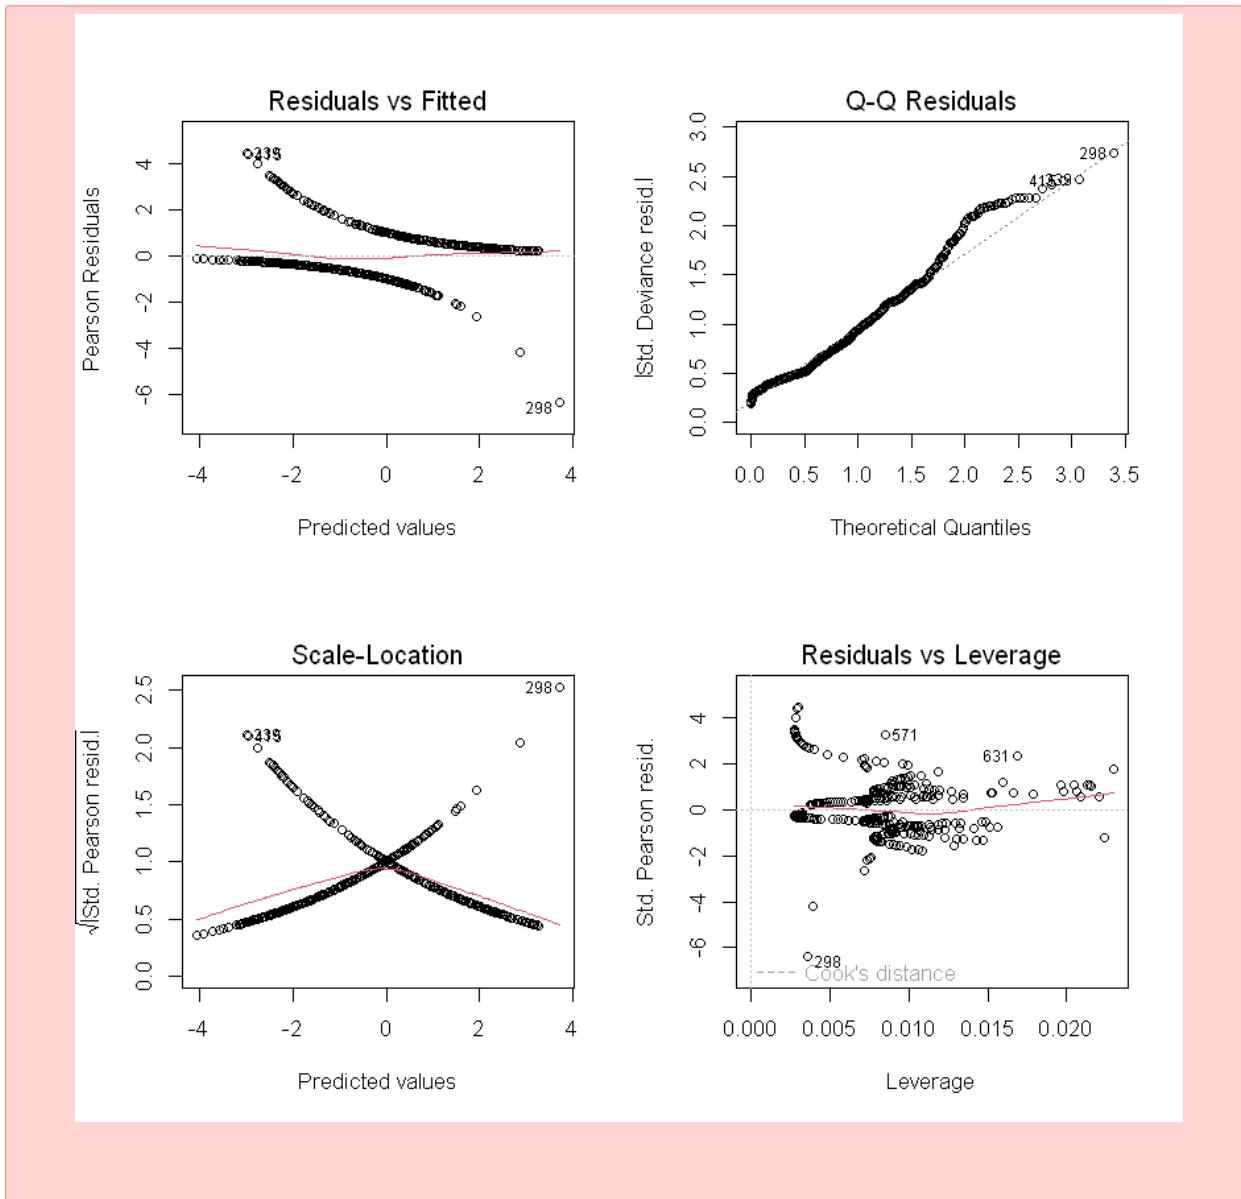
```
[7]: library (car)
print(vif(model))
```

	GVIF	Df	GVIF^(1/(2*Df))
Sex	1.072870	1	1.035794
Age	1.333112	1	1.154605
Pclass	1.416574	2	1.090962

I eventually ended up with logistic regression with the variables Sex, Age, and Pclass as my model. I first tried a decision tree and random forest approach but the accuracy was < 70% on the training sample. I also tried using K-Nearest neighbors, but it also didn't perform too well. For logistic regression, I decided to use Sex, Age, and Pclass because I thought they would be the most significant in predicting survival. I tried some other variables but they weren't as significant. All the predictors in this model are significant to 0.1% which is extremely good. Additionally all the VIFs (and adjusted VIFs) are < 2 indicating little to no multicollinearity.

Part 2: Residual Analysis

```
[8]: par(mfrow=c(2,2))
plot(model)
```



```
[9]: sresid = rstudent(model)
lev = hatvalues(model)
cook = cooks.distance(model)
```

```
[10]: print(sresid[which.max(abs(sresid))])
```

298
-2.756722

```
[11]: print(lev[which.max(abs(lev))])
```

```
484
0.02295272
```

```
[12]: print(cook[which.max(abs(cook))])
```

```
298
0.02904976
```

Looking at the graph, we see that the residuals match what we would expect from a logistic regression. The top band indicates a true of 1 which curves to 0 as the prediction goes from 0 to 1. Similarly, the bottom band indicates a true value of 0 which curves downward from 0 as the prediction increases. The QQ plot also is relatively linear indicating that the residuals are somewhat normally distributed (and deviations from normality is to be expected in logistic regression).

The largest studentized residual is -2.76 which is high but it is < 3 so it doesn't pose a huge problem.

The largest leverage is 0.023 which is slightly above the threshold of $3 * (\# \text{ of var})/n$ and is an outlier that could influence the fit.

The largest Cook's distance is 0.029 which is extremely small and means that the overall influence of all points is not significant.

The results suggest removing data points 298 and 484 to get a better model.

```
[13]: train = train[-c(298, 484),]
nrow(train)
model = glm(Survived~Sex+Age+Pclass, family="binomial", data=train)
summary(model)

glm.probs=predict(model,type="response")
glm.pred=rep(0,nrow(train))
glm.pred[glm.probs >0.457777]=1
table(glm.pred ,train$Survived)
mean(glm.pred==train$Survived)
```

```
712
```

```
Call:
```

```
glm(formula = Survived ~ Sex + Age + Pclass, family = "binomial",
     data = train)
```

```
Coefficients:
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	3.772411	0.400982	9.408	< 2e-16 ***
Sexmale	-2.518634	0.207387	-12.145	< 2e-16 ***
Age	-0.036952	0.007653	-4.828	1.38e-06 ***
Pclass2	-1.308652	0.277937	-4.708	2.50e-06 ***

```
Pclass3      -2.574697   0.281459  -9.148  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 962.43  on 711  degrees of freedom
Residual deviance: 646.89  on 707  degrees of freedom
AIC: 656.89

Number of Fisher Scoring iterations: 5

glm.pred  0   1
  0 348 62
  1 74 228
0.808988764044944
```

Part 3: Cross-validation

```
[14]: test = read.csv("TitanicTest.csv")
test = na.omit(test)
fix(test)
```

```
[15]: test$Sex = as.factor(test$Sex)
test$Pclass = as.factor(test$Pclass)
```

```
[16]: glm.test=predict(model,newdata=test,type="response")
glm.pred_test=rep(0,nrow(test))
glm.pred_test[glm.test >.5]=1

table(glm.pred_test,test$Survived)
mean(glm.pred_test==test$Survived)
```

```
glm.pred_test  0   1
  0 148 38
  1 45 88
0.739811912225705
```

The test model has around 74% accuracy which is pretty good!

Part (a)

```
[17]: test_resid = test$Survived - glm.test
mean(test_resid)
sd(test_resid)
```

```
-0.0408893896166924
0.411221458801901
```

[18]:

```
train_resid = train$Survived - glm.probs
mean(train_resid)
sd(train_resid)
```

```
-5.76050813939945e-16
0.378939961082023
```

The mean and standard deviation from the test set closely matches the corresponding values in the training set (mean is around 0 and std. dev is around 0.4)

Part (b)

[19]:

```
test_sresid = test_resid/sd(test_resid)
print(test_sresid[which.max(abs(test_sresid))])
```

```
261
2.27423
```

Computing the studentized residual we see that the largest value is 2.27 which is less than 3 so there are no outlier holdout cases.

Part (c)

[20]:

```
train_r2 = cor(glm.probs, train$Survived)^2
test_cor = cor(glm.test, test$Survived)
test_r2 = test_cor^2
shrinkage = train_r2 - test_r2

print(train_r2)
print(test_cor)
print(test_r2)
print(shrinkage)
```

```
[1] 0.4060841
[1] 0.5470174
[1] 0.299228
[1] 0.106856
```

The test correlation is 0.55 and the shrinkage is 0.11 which I think is small enough to say that the model predicts fairly well.

Part 4: Final Model

[21]: `summary(model)`

```

Call:
glm(formula = Survived ~ Sex + Age + Pclass, family = "binomial",
     data = train)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 3.772411   0.400982  9.408 < 2e-16 ***
Sexmale     -2.518634   0.207387 -12.145 < 2e-16 ***
Age        -0.036952   0.007653 -4.828 1.38e-06 ***
Pclass2     -1.308652   0.277937 -4.708 2.50e-06 ***
Pclass3     -2.574697   0.281459 -9.148 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 962.43 on 711 degrees of freedom
Residual deviance: 646.89 on 707 degrees of freedom
AIC: 656.89

Number of Fisher Scoring iterations: 5

```

[]: `pred = exp(3.772411 - 2.518634*Sexmale -0.036952*Age -1.308652*Pclass2 - 2.`
`-574697*Pclass3)/(1+exp(3.772411 - 2.518634*Sexmale -0.036952*Age -1.`
`-308652*Pclass2 - 2.574697*Pclass3))`

Sexmale is 1 if male and 0 if female. Pclass2 is 1 if Pclass is 2 and 0 otherwise and similarly, Pclass3 is 1 if Pclass is 3 and 0 otherwise:

[23]: `contrasts(train$Pclass)`

	2	3
1	0	0
2	1	0
3	0	1

Extra Credit: Precision and Recall

[24]: `table(glm.pred,train$Survived)`

glm.pred	0	1
----------	---	---

```
0 348 62  
1 74 228
```

```
[26]: train_precision = 348/(348+62)  
train_recall = 348/(348+74)  
  
print(train_precision)  
print(train_recall)
```

```
[1] 0.8487805  
[1] 0.8246445
```

```
[25]: table(glm.pred_test,test$Survived)
```

```
glm.pred_test 0 1  
0 148 38  
1 45 88
```

```
[27]: test_precision = 148/(148+38)  
test_recall = 148/(148+45)  
  
print(test_precision)  
print(test_recall)
```

```
[1] 0.7956989  
[1] 0.7668394
```

The precision and recall are both higher than the accuracy of the model and tell us the proportion of correct guess when the model predicts a death and how well a model predicts a death respectively.